



# Data Management and Artificial Intelligence

## Lecture 12

Sebastian Wandelt (小塞)

Beihang University


# Outline for today

- Review
- Multi-Armed Bandit Problem
- Monte Carlo Methods
- Monte-Carlo Tree Search (MCTS)

# Review

# Challenge: GA for the following problem

- Find maximum value of a function  $f(p)=31p - p^2$  with a single integer parameter  $p$  ( $0 \leq p \leq 31$ )




# Chromosome: Two options

1. An individual chromosome is a 1-integer number
  - $22=22$
2. An individual chromosome is a 5-bit number
  - $\mathbf{10110}=\mathbf{1}*2^4+\mathbf{0}*2^3+\mathbf{1}*2^2+\mathbf{1}*2^1+\mathbf{0}*2^0}=22$

# Chromosome example

- An individual chromosome is a 5-bit number
  - **10110** =  $1 \cdot 2^4 + 0 \cdot 2^3 + 1 \cdot 2^2 + 1 \cdot 2^1 + 0 \cdot 2^0 = 22$
- Four randomly generated genomes

Genome	p	Fitness
10110	22	198
00011	3	84
00010	2	58
11001	25	150



# Recombination

Genome	p	Fitness
<b>10110</b>	<b>22</b>	<b>198</b>
00011	3	84
00010	2	58
<b>11001</b>	<b>25</b>	<b>150</b>

- Recombination of 10110 and 11001 after bit 2:
  - Parent1: 10 110 (22)
  - Parent2: 11 001 (25)
- Offspring (before mutation):
  - Offspring1: **10001** (17) -> fitness=238
  - Offspring2: **11110** (30) -> fitness=30

# Summary: Components of a GA


A problem definition as input, and:

- Encoding principles (gene, chromosome)
- Initialization procedure (creation)
- Selection of parents (reproduction)
- Genetic operators (mutation, recombination)
- Evaluation function (environment)
- Termination condition



# Review

- What does the typical search process of the methods learned so far look like?



# Monte-Carlo Methods

# The core of Monte-Carlo

- Monte-Carlo methods are about randomness (again 😊)
- Before we get into Monte-Carlo, we have to discuss and understand randomness a bit more ...

# The core of Monte-Carlo

- Randomness
- Definition of random from Merriam-Webster:
  - Main Entry: **random**  
Function: *adjective*  
Date: 1565  
**1 a** : lacking a definite plan, purpose, or pattern **b** : made, done, or chosen at random <read *random* passages from the book>  
**2 a** : relating to, having, or being elements or events with definite probability of occurrence <*random* processes> **b** : being or relating to a set or to an element of a set each of whose elements has equal probability of occurrence <a *random* sample>; *also* : characterized by procedures designed to obtain such sets or elements <*random* sampling>

What is the obvious problem?



# Random Number

- What is random number?
- Is 3 a random number?
  - There is no such thing as single random number
- Random numbers
  - A sequence of numbers that have nothing to do with the other numbers in the sequence
    - 1,2,3,4,5,6,7,8,9,10,11,... ?
- In a uniform distribution of random numbers in the range  $[0,1]$ , every number has the same chance of turning up.
  - 0.00001 is just as likely as 0.5000

# Random v. Pseudo-random

- **Random numbers** have no defined sequence or formulation. Thus, for any  $n$  random numbers, each appears with equal probability.
- Computer algorithms are restricted to generating what we call **pseudo-random numbers**.
  - Generating a sequence of numbers whose properties approximate the properties of sequences of random numbers.
  - Based on one initial number only (the “seed”)

# Creating random numbers

- Any idea for how to create random numbers with a computer?
  - (“import random” is **not** the answer 😊)



# An early example (John von Neumann, 1946)

- To generate 10 digits of integer
  - Start with one of 10 digits integers
  - Square it and take middle 10 digits from answer
  - Example:  $5772156649^2 = 33317\underline{7923805949}09291$
- The sequence appears to be random, but each number is determined from the previous  $\rightarrow$  not random.
- Smaller example:
  - $6100^2 = 372\underline{10000}$
  - $2100^2 = 044\underline{10000}$
  - $4100^2 = 168\underline{10000}$
  - ...

Any problem?





# An early example (John Von Neumann, 1946)

- To generate 10 digits of integer
  - Start with one of 10 digits integers
  - Square it and take middle 10 digits from answer
  - Example:  $5772156649^2 = 33317\underline{7923805949}09291$
- The sequence appears to be random, but each number is determined from the previous → not random.
- Serious problem : Small numbers (0 or 1) are lumped together, it can get itself to a short loop. For example:
  - $6100^2 = 372\underline{10000}$
  - $2100^2 = 044\underline{10000}$
  - $4100^2 = 168\underline{10000}$
  - $8100^2 = 656\underline{10000}$
  - ☹

Initial number=**seed**

# RANDU Generator



- 1960's IBM
- Algorithm:

$$I_{n+1} = (65539 \times I_n) \bmod(2^{31})$$



Any problem?





# 1D and 2D Distribution of RANDU



# 3D Distribution of RANDU



Problems seen only when observed at the right angle



# There is a lot of research on that

- For now, just remember that pseudo-random numbers:
  1. Can be controlled by a seed
  2. Are in fact periodic sequences (e.g., when last number=seed)
  3. A finite set of quantized numbers -> problems in high-dimensions

# Initializing with Seeds

Two major reasons to initialize the seed:

1. The default state always generates the same sequence of random numbers. Not really random at all, particularly for a small set of calls. Solution: Call the seed method with the lower-order bits of the system clock.
2. You need a deterministic process that is repeatable.

Clear so far?



# A small intermediate challenge

- **Problem:** What is the probability that 10 dice throws add up exactly to 32?



# A small intermediate challenge

- **Problem:** What is the probability that 10 dice throws add up exactly to 32?
1. **Exact Way.** Calculate this exactly by counting all possible ways of making 32 from 10 dice.
  2. **Approximate (Lazy) Way.** Simulate throwing the dice (say 500 times), count the number of times the results add up to 32, and divide this by 500.



# Let us do this in Python



- **Problem:** What is the probability that 10 dice throws add up exactly to 32?
- **Solution:** Let's try the approximate way


# Source code for approximate simulation

```
import random
import matplotlib.pyplot as plt

fig, ax = plt.subplots(1, 1, figsize=(15, 10), dpi=70)
for N in range(1000, 150000, 1000):
    print(N)
    sum32 = 0
    for run in range(N):
        if sum([random.randint(1, 6) for i in range(10)]) == 32:
            sum32 += 1
    p = sum32 * 100 / N
    ax.plot(N, p, "kx")

ax.set_xlabel("N")
ax.set_ylabel("Probability")
```

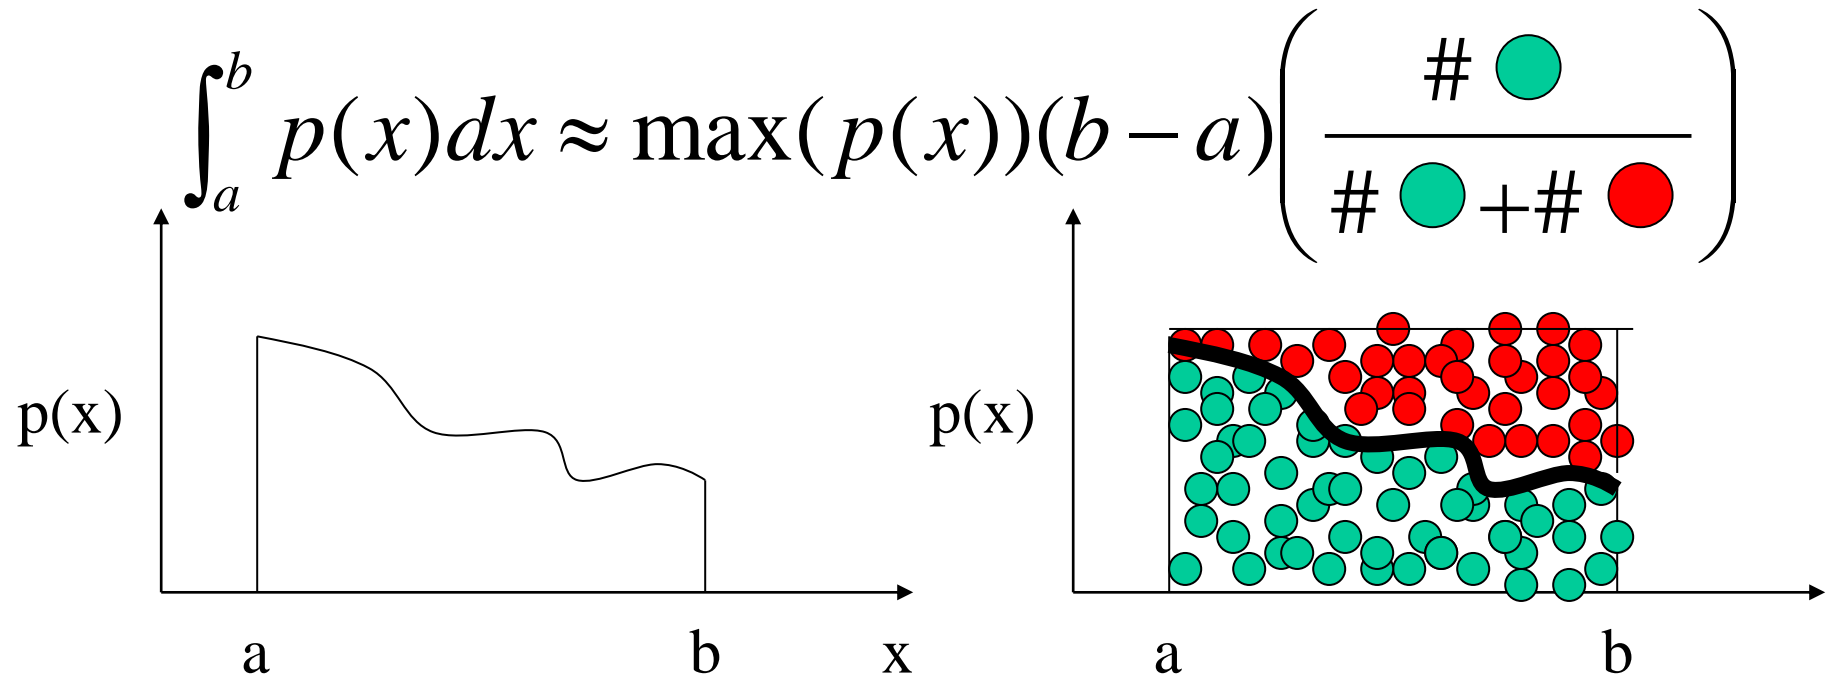
# Chart for approximate simulation



# Simple Example:

$$\int_a^b p(x) dx$$

- Method 1: Analytical Integration
- Method 2: Quadrature
- Method 3: MC -- random sampling the area enclosed by  $a < x < b$  and  $0 < y < \max(p(x))$



# Challenge: Estimating $\pi$ using Monte Carlo

- Can we estimate  $\pi$  using Monte-Carlo?



# Estimating $\pi$ using Monte Carlo


- The probability of a random point lying inside the unit circle:

$$P(x^2 + y^2 < 1) = \frac{A_{circle}}{A_{square}} = \frac{\pi}{4}$$

- If pick a random point  $N$  times and  $M$  of those times the point lies inside the unit circle:

$$P^{\circ}(x^2 + y^2 < 1) = \frac{M}{N}$$

- If  $N$  becomes very large,  $PI=P0$
- Let's try that in Python!



$$\pi = \frac{4 \cdot M}{N}$$



# Estimating $\pi$ using Monte Carlo

```
import random
import math
import matplotlib.pyplot as plt

inside,outside=[],[]
for i in range(10000000):
    x=random.uniform(-1,1)
    y=random.uniform(-1,1)
    if math.sqrt(x*x+y*y)<=1:
        inside.append((x,y))
    else:
        outside.append((x,y))

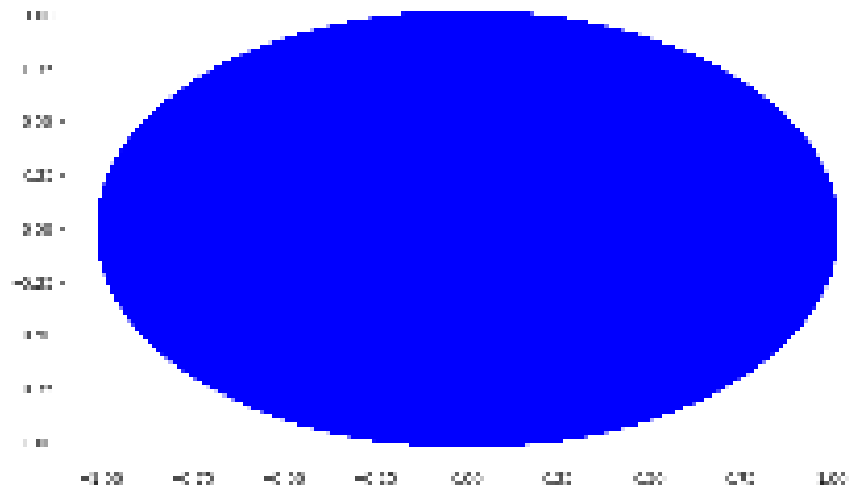
print(len(inside)*4/(len(outside)+len(inside)))
fig,ax=plt.subplots(1,1,figsize=(10,6),dpi=50)

xs,ys=list(zip(*inside))
ax.plot(xs,ys,"bo")
ax.plot([x for x,y in inside],[y for x,y in inside],"bo")
```

# Estimating $\pi$ using Monte Carlo

- Results:

– N =	10,000	Pi= 3.104385
– N =	100,000	Pi= 3.139545
– N =	1,000,000	Pi= 3.139668
– N =	10,000,000	Pi= 3.141774
– ...		





# Question:

- Do we need randomness here?



# Question:


- Do we need randomness here?
- In fact, this is really a sampling problem
  - We could also evenly distribute points along the plane and compute the ratio

# Multi-Armed Bandit Problem


# Single bandit



- Process:
  1. Pull arm of bandit
  2. Wait for three symbols to appear
    - If they are identical  $\Rightarrow$  win
    - If they are not identical  $\Rightarrow$  loose



# Multi-Armed Bandit Problem



**No two Slot Machines are the same!**

So: How to pick between Slot Machines so that you walk out with most \$\$\$ from Las Vegas?


# Problem




We are broke

Thousands of Slot Machines


We don't have enough budget  
to explore each Slot  
Machine even once!!



# An analogy in the “real world”



# Multi-Armed Bandit Problem




**Exploitation: Earn**

**Exploration: Learn**




# As a game?

Money = Computational Budget



Slot Machine = Next Action to Choose



Number of Slot Machines = Branching Factor

# UCB Algorithm for Minimizing Cumulative Regret


- $n(a)$  : number of pulls of arm  $a$  so far
- $Q(a)$  : average reward for trying action  $a$  so far
- Action choice by UCB after  $n$  pulls:

$$a_n = \arg \max_a Q(a) + \sqrt{\frac{2 \ln n}{n(a)}}$$

- Assumes rewards in  $[0,1]$ . We can always normalize given a bounded reward assumption

UCB= Upper Confidence Bound

# UCB: Bounded Sub-Optimality

$$a_n = \arg \max_a Q(a) + \sqrt{\frac{2 \ln n}{n(a)}}$$


## Value Term:

favors actions that looked good historically

## Exploration Term:

actions get an exploration bonus that grows with  $\ln(n)$

Doesn't waste much time on sub-optimal arms, unlike uniform!



# UCB Performance Guarantee

- Theorem: The expected cumulative regret of UCB  $E[Reg_n]$  after  $n$  arm pulls is bounded by  $O(\log n)$
- Is this good?
- Yes. The average per-step regret is  $O\left(\frac{\log(n)}{n}\right)$
- Theorem: No algorithm can achieve a better expected regret (up to constant factors)

# Task

- Let us implement the multi-armed bandit problem in Python!

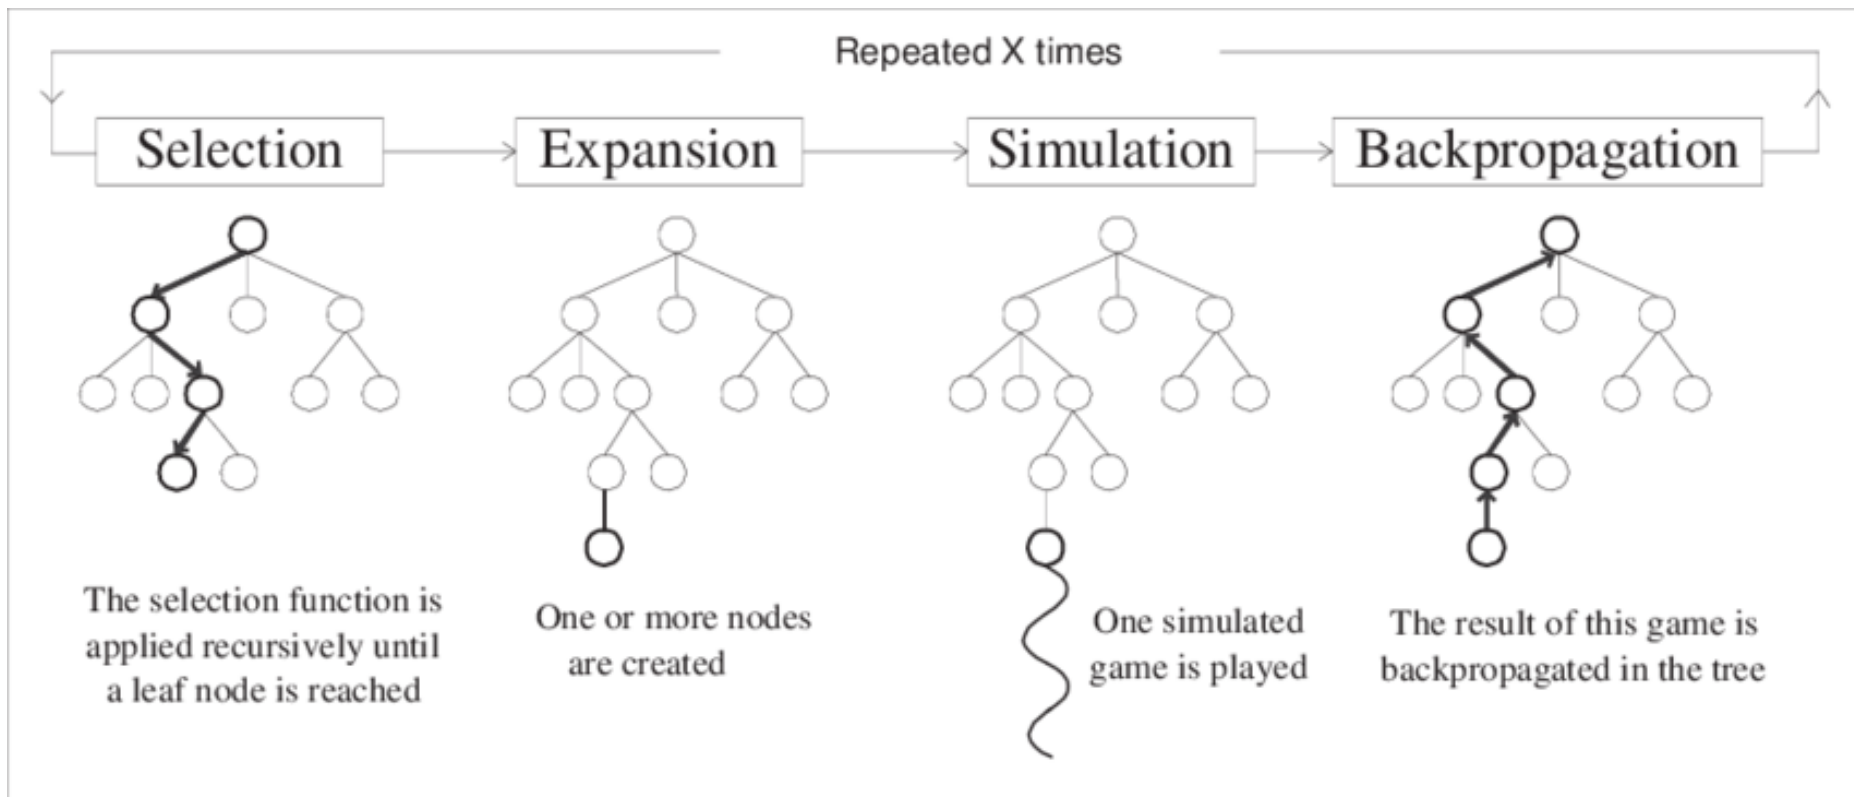


# **Generalization:**

## **Monte-Carlo Tree Search**

# Monte-Carlo Tree Search

- Builds a sparse look-ahead tree rooted at current state by repeated Monte-Carlo simulation of a “**rollout policy**”



# Summary

- Monte-Carlo algorithms and MCTS have revolutionized the area of game playing by computers
  - Alpha GO uses MCTS at its core
  - But the success of MCTS goes back to the early 2000s
- These algorithms have very strong theoretical properties (assuming infinite amount of time)
  - The more time you have (=the more you sample), the better results you get
- Many variants exist
  - Tuned for specific problem types



**Thank you very much!**